

# EDS 231 - Word Relationships

Juliet Cohen

2022-05-01

```
library(tidyr) #text analysis in R
library(pdftools)

## Using poppler version 22.02.0
library(lubridate) #working with date data

##
## Attaching package: 'lubridate'
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v dplyr   1.0.7
## v tibble  3.1.6      v stringr 1.4.0
## v readr   2.1.1      v forcats 0.5.1
## v purrr   0.3.4

## -- Conflicts ----- tidyverse_conflicts() --
## x lubridate::as.difftime() masks base::as.difftime()
## x lubridate::date()        masks base::date()
## x dplyr::filter()          masks stats::filter()
## x lubridate::intersect()   masks base::intersect()
## x dplyr::lag()              masks stats::lag()
## x lubridate::setdiff()     masks base::setdiff()
## x lubridate::union()       masks base::union()

library(tidytext)
library(readr)
library(quanteda)

## Package version: 3.2.1
## Unicode version: 13.0
## ICU version: 69.1

## Parallel computing: 8 of 8 threads used.

## See https://quanteda.io for tutorials and examples.
library(readtext) #quanteda subpackage for reading pdf
library(quanteda.textstats)
library(quanteda.textplots)
library(ggplot2)
```

```

library(forcats)
library(stringr)
library(quantda.textplots)
library(widyr) # pairwise correlations
library(igraph) # network plots

##
## Attaching package: 'igraph'

## The following object is masked from 'package:quantda.textplots':
##
##   as.igraph

## The following objects are masked from 'package:dplyr':
##
##   as_data_frame, groups, union

## The following objects are masked from 'package:purrr':
##
##   compose, simplify

## The following object is masked from 'package:tibble':
##
##   as_data_frame

## The following objects are masked from 'package:lubridate':
##
##   %--%, union

## The following object is masked from 'package:tidyr':
##
##   crossing

## The following objects are masked from 'package:stats':
##
##   decompose, spectrum

## The following object is masked from 'package:base':
##
##   union

library(ggraph)
library(patchwork)

```

## Import EPA EJ Data

```

setwd('.')
files <- list.files(pattern = "pdf$")

ej_reports <- lapply(files, pdf_text)

ej_pdf <- readtext("*.pdf", docvarsfrom = "filenames",
                  docvarnames = c("type", "subj", "year"),
                  sep = "_")

#creating an initial corpus containing our data

```

```
epa_corp <- corpus(x = ej_pdf, text_field = "text" )
summary(epa_corp)
```

```
## Corpus consisting of 6 documents, showing 6 documents:
```

```
##
##           Text Types Tokens Sentences type subj year
## EPA_EJ_2015.pdf 2136   8944         263 EPA   EJ 2015
## EPA_EJ_2016.pdf 1599   7965         176 EPA   EJ 2016
## EPA_EJ_2017.pdf 2774  16658         447 EPA   EJ 2017
## EPA_EJ_2018.pdf 3973  30564         653 EPA   EJ 2018
## EPA_EJ_2019.pdf 3773  22648         672 EPA   EJ 2019
## EPA_EJ_2020.pdf 4493  30523         987 EPA   EJ 2020
```

```
#I'm adding some additional, context-specific stop words to stop word lexicon
more_stops <-c("2015", "2016", "2017", "2018", "2019", "2020", "www.epa.gov", "https")
add_stops<- tibble(word = c(stop_words$word, more_stops))
stop_vec <- as_vector(add_stops)
```

Now we'll create some different data objects that will set us up for the subsequent analyses

```
#convert to tidy format and apply my stop words
raw_text <- tidy(epa_corp)
```

```
#Distribution of most frequent words across documents
```

```
raw_words <- raw_text %>%
  mutate(year = as.factor(year)) %>%
  unnest_tokens(word, text) %>%
  anti_join(add_stops, by = 'word') %>%
  count(year, word, sort = TRUE)
```

```
#number of total words by document
```

```
total_words <- raw_words %>%
  group_by(year) %>%
  summarize(total = sum(n))
```

```
report_words <- left_join(raw_words, total_words)
```

```
## Joining, by = "year"
```

```
# for the analysis that we want to do at the word level:
```

```
par_tokens <- unnest_tokens(raw_text, output = paragraphs, input = text, token = "paragraphs")
```

```
par_tokens <- par_tokens %>%
  mutate(par_id = 1:n())
```

```
par_words <- unnest_tokens(par_tokens, output = word, input = paragraphs, token = "words")
```

```
tokens <- tokens(epa_corp, remove_punct = TRUE) # create token obj
```

```
toks1<- tokens_select(tokens, min_nchar = 3)
```

```
toks1 <- tokens_tolower(toks1)
```

```
toks1 <- tokens_remove(toks1, pattern = (stop_vec)) # remove stop words
```

```
dfm <- dfm(toks1) # has docs in 1 col, the rows refer to num of occurrences for each word in the corpus
# fundamental obj for text analysis in quanteda
```

```
#first the basic frequency stat
```

```
tstat_freq <- textstat_frequency(dfm, n = 5, groups = year)
```

```
head(tstat_freq, 10)
```

```
##           feature frequency rank docfreq group
## 1  environmental      127    1      1  2015
## 2   communities      99    2      1  2015
## 3         epa        92    3      1  2015
## 4        justice      84    4      1  2015
## 5    community      47    5      1  2015
## 6  environmental     109    1      1  2016
## 7   communities      85    2      1  2016
## 8        justice      71    3      1  2016
## 9         epa        48    4      1  2016
## 10        federal      31    5      1  2016
```

1. What are the most frequent trigrams in the dataset? How does this compare to the most frequent bigrams? Which n-gram seems more informative here, and why?

Start by looking at bigrams:

```
toks2 <- tokens_ngrams(toks1, n=2) # bigram, tokenize, it goes thru the text with a 2 word window and c
dfm2 <- dfm(toks2)
dfm2 <- dfm_remove(dfm2, pattern = c(stop_vec))
freq_words2 <- textstat_frequency(dfm2, n=20)
freq_words2$token <- rep("bigram", 20)
#tokens1 <- tokens_select(tokens1, pattern = stopwords("en"), selection = "remove")

head(freq_words2)
```

```
##           feature frequency rank docfreq group token
## 1 environmental_justice      556    1      6  all bigram
## 2 technical_assistance      139    2      6  all bigram
## 3   drinking_water        133    3      6  all bigram
## 4   public_health        123    4      6  all bigram
## 5 progress_report        108    5      6  all bigram
## 6   air_quality          73    6      6  all bigram
```

The top 5 most frequent bigrams are:

1. environmental\_justice
2. technical\_assistance
3. drinking\_water
4. public\_health
5. progress\_report

```
toks2 <- tokens_ngrams(toks1, n = 3) # trigram, tokenize, it goes thru the text with a 3 word window an
dfm2 <- dfm(toks2)
dfm2 <- dfm_remove(dfm2, pattern = c(stop_vec))
freq_words2 <- textstat_frequency(dfm2, n=20)
freq_words2$token <- rep("trigram", 20)

head(freq_words2)
```

```
##           feature frequency rank docfreq group token
## 1 justice_fy2017_progress      51    1      1  all trigram
## 2 fy2017_progress_report      51    1      1  all trigram
```

|      |                                |    |   |   |             |
|------|--------------------------------|----|---|---|-------------|
| ## 3 | environmental_public_health    | 50 | 3 | 6 | all trigram |
| ## 4 | environmental_justice_fy2017   | 50 | 3 | 1 | all trigram |
| ## 5 | national_environmental_justice | 37 | 5 | 6 | all trigram |
| ## 6 | office_environmental_justice   | 32 | 6 | 6 | all trigram |

The top 5 most frequent trigrams are:

1. justice\_fy2017\_progress
2. fy2017\_progress\_report
3. environmental\_public\_health
4. environmental\_justice\_fy2017
5. national\_environmental\_justice

The trigrams show more repetitive words such as justice, progress, fy2017, and environmental, and appear to be words that do not form a sensical, stand-alone phrase when read together, while the bigrams are more diverse and the words make sense when read together in sequence. Therefore I think that bigrams are more informative here.

**2. Choose a new focal term to replace “justice” and recreate the correlation table and network (see corr\_paragraphs and corr\_network chunks). Explore some of the plotting parameters in the cor\_network chunk to see if you can improve the clarity or amount of information your plot conveys. Make sure to use a different color for the ties!**

```
word_cors <- par_words %>%
  add_count(par_id) %>%
  filter(n >= 50) %>%
  select(-n) %>%
  pairwise_cor(word, par_id, sort = TRUE) # generates correlation coefficients rather than just the num
# cols = item1 and item2 and correlation

wildlife_cors <- word_cors %>%
  filter(item1 == "wildlife")

word_cors %>%
  filter(item1 %in% c("wildlife", "environmental", "equity", "income")) %>%
  group_by(item1) %>%
  top_n(6) %>%
  #slice_max(item1, n = 6) %>%
  ungroup() %>%
  mutate(item1 = as.factor(item1),
         name = reorder_within(item2, correlation, item1)) %>%
  ggplot(aes(y = name, x = correlation, fill = item1)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~item1, ncol = 2, scales = "free")+
  scale_y_reordered() +
  labs(y = NULL,
       x = NULL,
       title = "Correlations with key words",
       subtitle = "EPA EJ Reports")
```

```
## Selecting by correlation
```

```
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'epa's' in 'mbcsToSbcs': dot substituted for <e2>
```





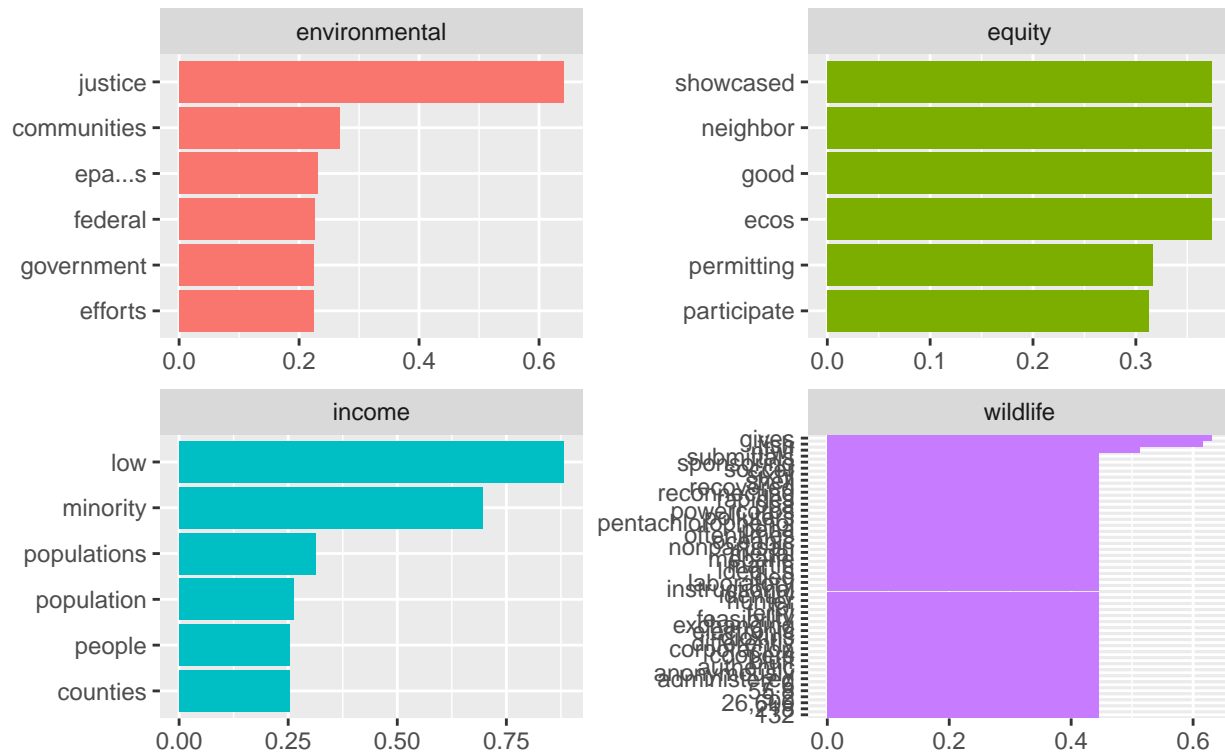
[illegible]



[illegible]

## Correlations with key words

### EPA EJ Reports



*#let's zoom in on just one of our key terms*

```
wildlife_cors <- word_cors %>%
  filter(item1 == "wildlife") %>%
  mutate(n = 1:n())
```

```
head(wildlife_cors)
```

```
## # A tibble: 6 x 4
##   item1   item2 correlation    n
##   <chr>   <chr>      <dbl> <int>
## 1 wildlife gives      0.631     1
## 2 wildlife fish       0.616     2
## 3 wildlife nfwf       0.513     3
## 4 wildlife 218       0.446     4
## 5 wildlife 132       0.446     5
## 6 wildlife 55.8      0.446     6
```

Not surprisingly, the correlation between “environmental” and “justice” is by far the highest, which makes sense given the nature of these reports. How might we visualize these correlations to develop of sense of the context in which justice is discussed here?

```
wildlife_cors %>%
  filter(n <= 50) %>%
  graph_from_data_frame() %>%
  ggraph(layout = "fr") +
  geom_edge_link(aes(edge_alpha = correlation, edge_width = correlation), edge_colour = "firebrick") +
  geom_node_point(size = 5) +
  geom_node_text(aes(label = name), repel = TRUE,
```

```

        point.padding = unit(0.2, "lines")) +
theme_void()

```

```

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>

```

```
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <99>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <e2>  
  
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, :  
## conversion failure on 'ldeq's' in 'mbcsToSbcs': dot substituted for <80>
```

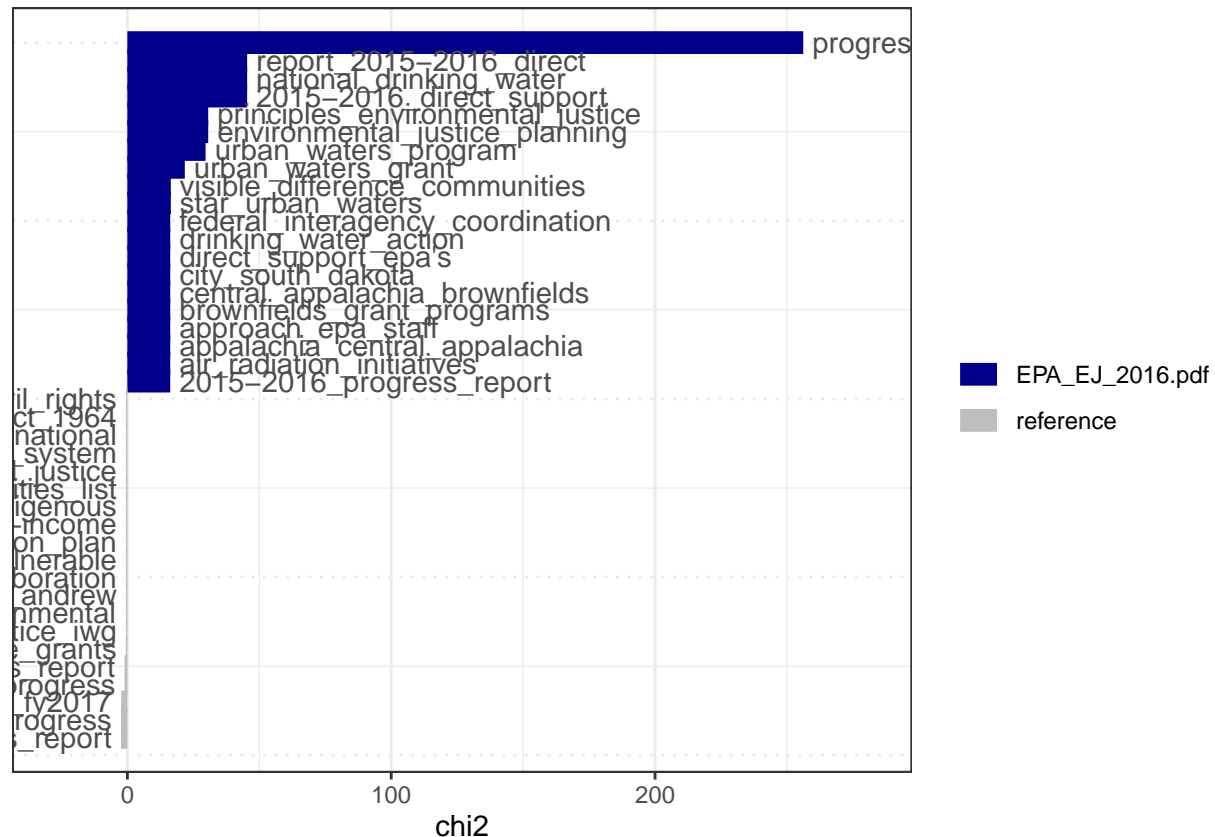


3. Write a function that allows you to conduct a keyness analysis to compare two individual EPA reports (hint: that means target and reference need to both be individual reports). Run the function on 3 pairs of reports, generating 3 keyness plots.

Use all of the frequencies for each word in each document and calculate a chi-square to see which words occur significantly more or less within a particular target document.

Example from class:

```
keyness <- textstat_keyness(dfm2, target = 2)
textplot_keyness(keyness)
```

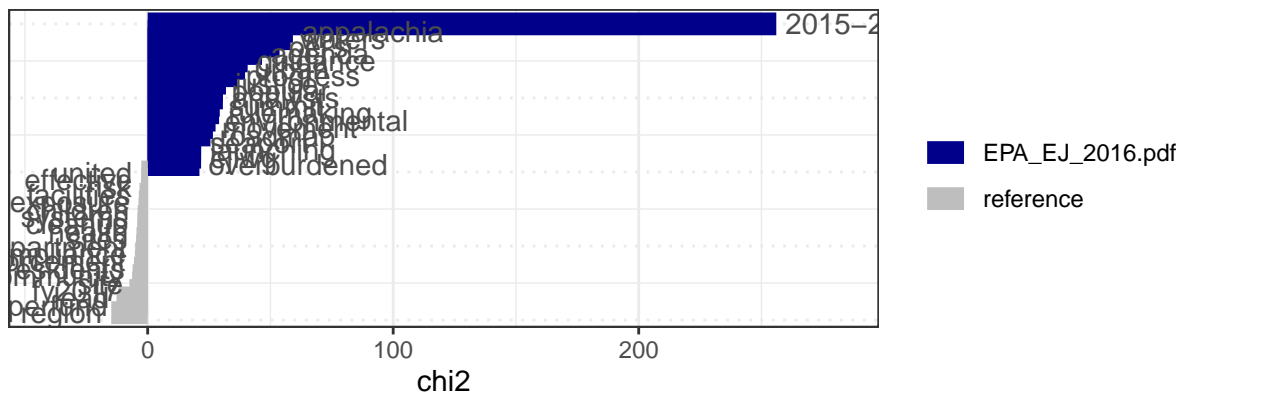
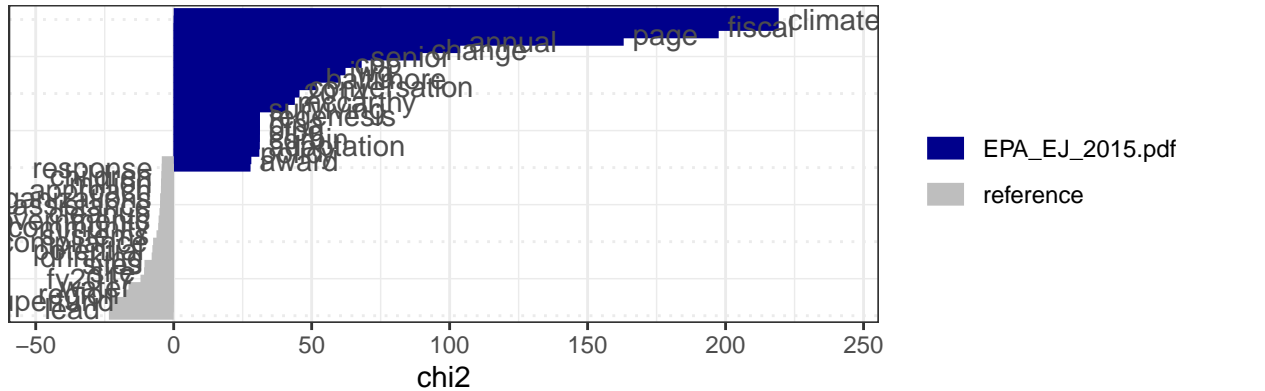


Create function:

```
keyness_analysis <- function(data, target1, target2) {
  keyness1 <- textstat_keyness(data, target = target1)
  keyness2 <- textstat_keyness(data, target = target2)
  plot1 <- textplot_keyness(keyness1)
  plot2 <- textplot_keyness(keyness2)
  plot <- plot1 / plot2
  return(plot)
}
```

Test Function:

```
keyness_analysis(data = dfm, target1 = 1, target2 = 2)
```



Select a word or multi-word term of interest and identify words related to it using windowing and keyness comparison. To do this you will create two objects: one containing all words occurring within a 10-word window of your term of interest, and the second object containing all other words. Then run a keyness comparison on these objects. Which one is the target, and which the reference? Hint